

---

# pyDive Documentation

*Release 1.1 beta*

**Heiko Bureau**

March 29, 2015



## CONTENTS

<b>1</b>	<b>Getting started</b>	<b>3</b>
1.1	Quickstart . . . . .	3
1.2	Setup an IPython.parallel cluster configuration . . . . .	3
1.3	Run tests . . . . .	4
1.4	Overview . . . . .	4
<b>2</b>	<b>Tutorials</b>	<b>5</b>
2.1	Example 1: Total field energy . . . . .	6
2.2	Example 2: Particle density field . . . . .	7
2.3	Example 3: Particle energy spectrum . . . . .	9
<b>3</b>	<b>Reference</b>	<b>11</b>
3.1	Arrays . . . . .	11
3.2	Modules . . . . .	18
<b>4</b>	<b>Indices and tables</b>	<b>29</b>
	<b>Python Module Index</b>	<b>31</b>
	<b>Index</b>	<b>33</b>



Contents:



## GETTING STARTED

### 1.1 Quickstart

pyDive is built on top of *IPython.parallel*, *numpy*, *mpi4py* and *h5py*. Running `python setup.py install` will install pyDive with these and other required packages from *requirements.txt*. Alternatively you can install it via pip: `pip install pyDive`.

Basic code example:

```
import pyDive
pyDive.init()

arrayA = pyDive.ones([1000, 1000, 1000], distaxis=0)
arrayB = pyDive.zeros_like(arrayA)

# do some array operations, + - * / sin cos, ..., slicing, etc...
...

# get numpy-array
result = arrayC.gather()
# plot result
...
```

Before actually running this script there must have been an *IPython.parallel* cluster launched (see section below) otherwise `pyDive.init()` fails.

To keep things simple pyDive distributes array-memory only along **one** user-specified axis. This axis is given by the *distaxis* parameter at array instantiation. It should usually be the largest axis in order to have the best surface-to-volume ratio. But keep in mind that during arithmetic operations both arrays have to be distributed along the *same* axis.

Although the array elements are stored on the cluster nodes you have full access through indexing. If you want to have a numpy-array from a pyDive-array anyway you can call the method `arrayC.gather()` but make sure that your pyDive-array is small enough to fit into your local machine's memory. If not you may want to slice it first.

### 1.2 Setup an IPython.parallel cluster configuration

The first step is to create an *IPython.parallel* profile: [http://ipython.org/ipython-doc/2/parallel/parallel\\_process.html](http://ipython.org/ipython-doc/2/parallel/parallel_process.html). The name of this profile is the argument of `pyDive.init()`. It defaults to "mpi". Starting the cluster is then the second and final step:

```
$ ipcluster start -n 4 --profile=mpi
```

## 1.3 Run tests

In order to test the pyDive installation you can run:

```
$ python setup.py test
```

This will ask you for the IPython.parallel profile to be used and the number of engines to be started, e.g.:

```
$ Name of your IPython-parallel profile you want to run the tests with: pbs
$ Number of engines: 256
```

Then the script starts the cluster, runs the tests and finally stops the cluster. If you have already a cluster running by your own you can also run the tests by launching `py.test` from the pyDive directory and setting the environment variable `IPP_PROFILE_NAME` to the profile's name.

## 1.4 Overview

**pyDive knows different kinds of distributed arrays, all corresponding to a local, non-distributed array.**

- `numpy` -> `pyDive.ndarray` -> Stores array elements in cluster nodes' memory.
- `hdf5` -> `pyDive.arrays.h5_ndarray` -> Stores array elements in a hdf5-file.
- `adios` -> `pyDive.arrays.ad_ndarray` -> Stores array elements in a adios-file.
- `pyDive.cloned_ndarray` -> Holds independent copies of one array on cluster nodes.

**Among these three packages there are a few modules:**

- `pyDive.arrayOfStructs` -> structured datatypes
- `pyDive.algorithm` -> map, reduce, mapReduce
- `pyDive.fragment` -> fragment file-disk array to fit into the cluster's main memory
- `pyDive.mappings` -> particle-mesh mappings
- `pyDive.picongpu` -> helper functions for picongpu-users
- `pyDive.pyDive` -> shortcuts for most used functions



## TUTORIALS

In this section we are going through a few use cases for pyDive. If you want to test the code you can download the sample hdf5-file. It has the following dataset structure:

```
$ h5ls -r sample.h5
/                               Group
/fields                         Group
/fields/fieldB                  Group
/fields/fieldB/z                Dataset {256, 256}
/fields/fieldE                  Group
/fields/fieldE/x                Dataset {256, 256}
/fields/fieldE/y                Dataset {256, 256}
/particles                      Group
/particles/cellidx              Group
/particles/cellidx/x            Dataset {10000}
/particles/cellidx/y            Dataset {10000}
/particles/pos                  Group
/particles/pos/x                Dataset {10000}
/particles/pos/y                Dataset {10000}
/particles/vel                  Group
/particles/vel/x                Dataset {10000}
/particles/vel/y                Dataset {10000}
/particles/vel/z                Dataset {10000}
```

After launching the cluster (*Setup an IPython.parallel cluster configuration*) the first step is to initialize pyDive:

```
import pyDive
pyDive.init()
```

Load a single dataset:

```
h5fieldB_z = pyDive.h5.fromPath("sample.h5", "/fields/fieldB/z", distaxis=0)

assert type(h5fieldB_z) is pyDive.h5.h5_ndarray
```

*h5fieldB\_z* just holds a dataset *handle*. To read out data into memory call `load()`:

```
fieldB_z = h5fieldB_z.load()

assert type(fieldB_z) is pyDive.ndarray
```

This loads the entire dataset into the main memory of all *engines*. The array elements are distributed along `distaxis=0`.

We can also load a hdf5-group:

```
h5fieldE = pyDive.h5.open("sample.h5", "/fields/fieldE", distaxis=0)
fieldE = h5fieldE.load()
```

*h5fieldE* and *fieldE* are some so called “virtual array-of-structures”, see: *pyDive.arrayOfStructs*.

```
>>> print h5fieldE
VirtualArrayOfStructs<array-type: <class 'pyDive.distribution.single_axis.h5_ndarray'>, shape: [256, 256],
  y -> float32
  x -> float32

>>> print fieldE
VirtualArrayOfStructs<array-type: <class 'pyDive.distribution.single_axis.ndarray'>, shape: [256, 256],
  y -> float32
  x -> float32
```

Now, let’s do some calculations!

## 2.1 Example 1: Total field energy

Computing the total field energy of an electromagnetic field means squaring and summing or in pyDive’s words:

```
import pyDive
import numpy as np
pyDive.init()

h5input = "sample.h5"

h5fields = pyDive.h5.open(h5input, "/fields", distaxis=0)
fields = h5fields.load() # read out all fields into cluster's main memory in parallel

energy_field = fields.fieldE.x**2 + fields.fieldE.y**2 + fields.fieldB.z**2

total_energy = pyDive.reduce(energy_field, np.add)
print total_energy
```

Output:

```
$ python example1.py
557502.0
```

Well this was just a very small hdf5-sample of 1.3 MB however in real world we deal with a lot greater data volumes. So what happens if *h5fields* is too large to be stored in the main memory of the whole cluster? The line `fields = h5fields.load()` will crash. In this case we want to load the hdf5 data piece by piece. The function *pyDive.fragment* helps us doing so:

```
import pyDive
import numpy as np
pyDive.init()

h5input = "sample.h5"

big_h5fields = pyDive.h5.open(h5input, "/fields", distaxis=0)
# big_h5fields.load() # would cause a crash

total_energy = 0.0
for h5fields in pyDive.fragment(big_h5fields):
    fields = h5fields.load()
```

```

energy_field = fields.fieldE.x**2 + fields.fieldE.y**2 + fields.fieldB.z**2

total_energy += pyDive.reduce(energy_field, np.add)

print total_energy

```

An equivalent way to get this result is a `pyDive.mapReduce`:

```

...
def square_fields(h5fields):
    fields = h5fields.load()
    return fields.fieldE.x**2 + fields.fieldE.y**2 + fields.fieldB.z**2

total_energy = pyDive.mapReduce(square_fields, np.add, h5fields)
print total_energy

```

`square_fields` is called on each *engine* where `h5fields` is a structure (`pyDive.arrayOfStructs`) of `h5_ndarrays` representing a sub part of the big `h5fields`. `pyDive.algorithm.mapReduce()` can be called with an arbitrary number of arrays including `pyDive.ndarrays`, `pyDive.h5.h5_ndarrays`, `pyDive.adios.ad_ndarrays` and `pyDive.cloned_ndarrays`. If there are `pyDive.h5.h5_ndarrays` or `pyDive.adios.ad_ndarrays` it will check whether they fit into the combined main memory of all cluster nodes as a whole and loads them piece by piece if not.

Now let's say our dataset is really big and we just want to get a first estimate of the total energy:

```

...
total_energy = pyDive.mapReduce(square_fields, np.add, h5fields[:, :10, ::10]) * 10.0**2

```

Slicing on `pyDive`-arrays is always allowed.

If you use `picongpu` here is an example of how to get the total field energy for each timestep (see `pyDive.picongpu`):

```

import pyDive
import numpy as np
pyDive.init()

def square_field(h5field):
    field = h5field.load()
    return field.x**2 + field.x**2 + field.x**2

for step, h5field in pyDive.picongpu.loadAllSteps("../simOutput", "fields/FieldE", distaxis=0):
    total_energy = pyDive.mapReduce(square_field, np.add, h5field)

    print step, total_energy

```

## 2.2 Example 2: Particle density field

Given the list of particles in our `sample.h5` we want to create a 2D density field out of it. For this particle-to-mesh mapping we need to apply a certain particle shape like cloud-in-cell (CIC), triangular-shaped-cloud (TSC), and so on. A list of these together with the actual mapping functions can be found in the `pyDive.mappings` module. If you miss a shape you can easily create one by your own by basically defining a particle shape function. Note that if you have `numba` installed the shape function will be compiled resulting in a significant speed-up.

We assume that the particle positions are distributed randomly. This means although each engine is loading a separate part of all particles it needs to write to the entire density field. Therefore the

density field must have a whole representation on each participating engine. This is the job of `pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray`.

```
import pyDive
import numpy as np
pyDive.init()

shape = [256, 256]
density = pyDive.cloned.zeros(shape)

h5input = "sample.h5"

particles = pyDive.h5.fromPath(h5input, "/particles", distaxis=0)

def particles2density(particles, density):
    particles = particles.load()
    total_pos = particles.cellidx.astype(np.float32) + particles.pos

    # convert total_pos to an (N, 2) shaped array
    total_pos = np.hstack((total_pos.x[:,np.newaxis],
                           total_pos.y[:,np.newaxis]))


    par_weighting = np.ones(particles.shape)
    import pyDive.mappings
    pyDive.mappings.particles2mesh(density, par_weighting, total_pos, pyDive.mappings.CIC)

pyDive.map(particles2density, particles, density)

final_density = density.sum() # add up all local copies

from matplotlib import pyplot as plt
plt.imshow(final_density)
plt.show()
```

Output:



Here, as in the first example, `particles2density` is a function executed on the *engines* by `pyDive.algorithm.map()`. All of its arguments are numpy-arrays or structures (`pyDive.arrayOfStructs`) of numpy-arrays.

`pyDive.algorithm.map()` can also be used as a decorator:

```
@pyDive.map
def particles2density(particles, density):
    ...

particles2density(particles, density)
```

## 2.3 Example 3: Particle energy spectrum

```
import pyDive
import numpy as np
pyDive.init()

bins = 256
spectrum = pyDive.cloned.zeros([bins])

h5input = "sample.h5"

velocities = pyDive.h5.fromPath(h5input, "/particles/vel", distaxis=0)
```

```
@pyDive.map
def vel2spectrum(velocities, spectrum, bins):
    velocities = velocities.load()
    mass = 1.0
    energies = 0.5 * mass * (velocities.x**2 + velocities.y**2 + velocities.z**2)


    spectrum[:, bin_edges] = np.histogram(energies, bins)

vel2spectrum(velocities, spectrum, bins=bins)

final_spectrum = spectrum.sum() # add up all local copies

from matplotlib import pyplot as plt
plt.plot(final_spectrum)
plt.show()
```

Output:



## REFERENCE

### 3.1 Arrays

#### 3.1.1 pyDive.arrays.ndarray module

---

**Note:** All of this module's functions and classes are also directly accessible from the `pyDive` module.

---

##### `pyDive.ndarray` class

**class** `pyDive.ndarray` (*shape*, *dtype*=<type 'float'>, *distaxis*=0, *target\_offsets*=None, *target\_ranks*=None, *no\_allocation*=False, *\*\*kwargs*)

Represents a cluster-wide, multidimensional, homogeneous array of fixed-size elements. *cluster-wide* means that its elements are distributed across IPython.parallel-engines. The distribution is done in one dimension along a single, user-specified axis. The user can optionally specify which engine maps to which index range or leave the default that pursues an uniform distribution across all engines.

This `ndarray` - class is auto-generated out of its local counterpart: `numpy.ndarray`.

The implementation is based on IPython.parallel and local `numpy.ndarray` - arrays. Every special operation `numpy.ndarray` implements (“\_\_add\_\_”, “\_\_le\_\_”, ...) is also available for `ndarray`.

Note that array slicing is a cheap operation since no memory is copied. However this can easily lead to the situation where you end up with two arrays of the same size but of distinct element distribution. Therefore call `dist_like()` first before doing any manual stuff on their local arrays. However every cluster-wide array operation first equalizes the distribution of all involved arrays, so an explicit call to `dist_like()` is rather unlikely in most use cases.

**\_\_init\_\_** (*shape*, *dtype*=<type 'float'>, *distaxis*=0, *target\_offsets*=None, *target\_ranks*=None, *no\_allocation*=False, *\*\*kwargs*)

Creates an instance of `ndarray`. This is a low-level method of instantiating an array, it should rather be constructed using factory functions (“empty”, “zeros”, “open”, ...)

##### Parameters

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **target\_offsets** (*ints*) – list of indices marking the offset along the distributed axis of each local array.
- **target\_ranks** (*ints*) – list of *engine* ranks holding the local arrays.

- **no\_allocation** (*bool*) – if `True` no instance of `numpy.ndarray` will be created on engine. Useful for manual instantiation of the local array.
- **kwargs** – additional keyword arguments are forwarded to the constructor of the local array.

**copy** ()

Returns a hard copy of this array.

**dist\_like** (*other*)

Redistributes a copy of this array (*self*) like *other* and returns the result. Checks whether redistribution is necessary and returns *self* if not.

Redistribution involves inter-engine communication.

**Parameters** *other* (*distributed array*) – target array

**Raises**

- **AssertionError** – if the shapes of *self* and *other* don't match.
- **AssertionError** – if *self* and *other* are distributed along distinct axes.

**Returns** new array with the same content as *self* but distributed like *other*. If *self* is already distributed like *other* nothing is done and *self* is returned.

**gather** ()

Gathers local instances of `numpy.ndarray` from *engines*, concatenates them and returns the result.

**Returns** instance of `numpy.ndarray`

## Factory functions

These are convenient functions to create a *pyDive.ndarray* instance.

`pyDive.arrays.ndarray.array` (*array\_like*, *distaxis=0*)

Create a *pyDive.ndarray* instance from an array-like object.

**Parameters**

- **array\_like** – Any object exposing the array interface, e.g. `numpy-array`, `python sequence`, ...
- **distaxis** (*int*) – distributed axis

`pyDive.arrays.ndarray.empty` (*shape*, *dtype=<type 'float'>*, *distaxis=0*, *\*\*kwargs*)

Create a *ndarray* instance. This function calls its local counterpart *numpy.empty* on each *engine*.

**Parameters**

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **kwargs** – keyword arguments are passed to the local function *numpy.empty*

`pyDive.arrays.ndarray.empty_like` (*other*, *\*\*kwargs*)

Create a *ndarray* instance with the same shape, dtype and distribution as *other*. This function calls its local counterpart *numpy.empty\_like* on each *engine*.

**Parameters**

- **other** – other array



- **kwargs** – keyword arguments are passed to the local function *numpy.empty\_like*

`pyDive.arrays.ndarray.hollow(shape, dtype=<type 'float'>, distaxis=0)`

Create a `pyDive.ndarray` instance distributed across all engines without allocating a local `numpy`-array.

#### Parameters

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis

`pyDive.arrays.ndarray.hollow_like(other)`

Create a `pyDive.ndarray` instance with the same shape, distribution and type as `other` without allocating a local `numpy`-array.

`pyDive.arrays.ndarray.zeros(shape, dtype=<type 'float'>, distaxis=0, **kwargs)`

Create a `ndarray` instance. This function calls its local counterpart *numpy.zeros* on each *engine*.

#### Parameters

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **kwargs** – keyword arguments are passed to the local function *numpy.zeros*

`pyDive.arrays.ndarray.zeros_like(other, **kwargs)`

Create a `ndarray` instance with the same shape, dtype and distribution as `other`. This function calls its local counterpart *numpy.zeros\_like* on each *engine*.

#### Parameters

- **other** – other array
- **kwargs** – keyword arguments are passed to the local function *numpy.zeros\_like*

`pyDive.arrays.ndarray.ones(shape, dtype=<type 'float'>, distaxis=0, **kwargs)`

Create a `ndarray` instance. This function calls its local counterpart *numpy.ones* on each *engine*.

#### Parameters

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **kwargs** – keyword arguments are passed to the local function *numpy.ones*

`pyDive.arrays.ndarray.ones_like(other, **kwargs)`

Create a `ndarray` instance with the same shape, dtype and distribution as `other`. This function calls its local counterpart *numpy.ones\_like* on each *engine*.

#### Parameters

- **other** – other array
- **kwargs** – keyword arguments are passed to the local function *numpy.ones\_like*

## Universal functions

*numpy* knows the so called *ufuncs* (universal function). These are functions which can be applied elementwise on an array, like *sin*, *cos*, *exp*, *sqrt*, etc. All of these *ufuncs* from *numpy* are also available for *pyDive.ndarray* arrays, e.g.

```
a = pyDive.ones([100])
a = pyDive.sin(a)
```

### 3.1.2 pyDive.arrays.h5\_ndarray module

---

**Note:** This module has a shortcut: `pyDive.h5`.

---

**class** `pyDive.arrays.h5_ndarray.h5_ndarray` (*shape*, *dtype*=<type 'float'>, *distaxis*=0, *target\_offsets*=None, *target\_ranks*=None, *no\_allocation*=False, *\*\*kwargs*)

Represents a cluster-wide, multidimensional, homogeneous array of fixed-size elements. *cluster-wide* means that its elements are distributed across IPython.parallel-engines. The distribution is done in one dimension along a single, user-specified axis. The user can optionally specify which engine maps to which index range or leave the default that pursues an uniform distribution across all engines.

This **h5\_ndarray** - class is auto-generated out of its local counterpart: **pyDive.arrays.h5\_ndarray.h5\_ndarray\_local**.

The implementation is based on IPython.parallel and local `pyDive.arrays.h5_ndarray.h5_ndarray_local` - arrays. Every special operation `pyDive.arrays.h5_ndarray.h5_ndarray_local` implements (“\_\_add\_\_”, “\_\_le\_\_”, ...) is also available for `h5_ndarray`.

Note that array slicing is a cheap operation since no memory is copied. However this can easily lead to the situation where you end up with two arrays of the same size but of distinct element distribution. Therefore call `dist_like()` first before doing any manual stuff on their local arrays. However every cluster-wide array operation first equalizes the distribution of all involved arrays, so an explicit call to `dist_like()` is rather unlikely in most use cases.

**\_\_init\_\_** (*shape*, *dtype*=<type 'float'>, *distaxis*=0, *target\_offsets*=None, *target\_ranks*=None, *no\_allocation*=False, *\*\*kwargs*)

Creates an instance of `h5_ndarray`. This is a low-level method of instantiating an array, it should rather be constructed using factory functions (“empty”, “zeros”, “open”, ...)

#### Parameters

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **target\_offsets** (*ints*) – list of indices marking the offset along the distributed axis of each local array.
- **target\_ranks** (*ints*) – list of *engine* ranks holding the local arrays.
- **no\_allocation** (*bool*) – if True no instance of `pyDive.arrays.h5_ndarray.h5_ndarray_local` will be created on engine. Useful for manual instantiation of the local array.
- **kwargs** – additional keyword arguments are forwarded to the constructor of the local array.

**load()**

Load array from file into main memory of all engines in parallel.

**Returns** pyDive.ndarray instance

`pyDive.arrays.h5_ndarray.open(filename, datapath, distaxis=0)`

Create an pyDive.h5.h5\_ndarray instance respectively a structure of pyDive.h5.h5\_ndarray instances from file.

**Parameters**

- **filename** – name of hdf5 file.
- **dataset\_path** – path within hdf5 file to a single dataset or hdf5 group.
- **int** (*distaxis*) – distributed axis

**Returns** pyDive.h5.h5\_ndarray instance / structure of pyDive.h5.h5\_ndarray instances

`pyDive.arrays.h5_ndarray.open_dset(filename, dataset_path, distaxis=0)`

Create a pyDive.h5.h5\_ndarray instance from file.

**Parameters**

- **filename** – name of hdf5 file.
- **dataset\_path** – path within hdf5 file to a single dataset.
- **int** (*distaxis*) – distributed axis

**Returns** pyDive.h5.h5\_ndarray instance

### 3.1.3 pyDive.arrays.ad\_ndarray module

---

**Note:** This module has a shortcut: `pyDive.adios`.

---

```
class pyDive.arrays.ad_ndarray.ad_ndarray(shape, dtype=<type 'float'>, distaxis=0,
                                          target_offsets=None, target_ranks=None,
                                          no_allocation=False, **kwargs)
```

Represents a cluster-wide, multidimensional, homogeneous array of fixed-size elements. *cluster-wide* means that its elements are distributed across IPython.parallel-engines. The distribution is done in one dimension along a single, user-specified axis. The user can optionally specify which engine maps to which index range or leave the default that pursues an uniform distribution across all engines.

This **ad\_ndarray** - class is auto-generated out of its local counterpart: **pyDive.arrays.ad\_ndarray.ad\_ndarray\_local**.

The implementation is based on IPython.parallel and local `pyDive.arrays.ad_ndarray.ad_ndarray_local` - arrays. Every special operation `pyDive.arrays.ad_ndarray.ad_ndarray_local` implements (“\_\_add\_\_”, “\_\_le\_\_”, ...) is also available for `ad_ndarray`.

Note that array slicing is a cheap operation since no memory is copied. However this can easily lead to the situation where you end up with two arrays of the same size but of distinct element distribution. Therefore call `dist_like()` first before doing any manual stuff on their local arrays. However every cluster-wide array operation first equalizes the distribution of all involved arrays, so an explicit call to `dist_like()` is rather unlikely in most use cases.

```
__init__(shape, dtype=<type 'float'>, distaxis=0, target_offsets=None, target_ranks=None,
         no_allocation=False, **kwargs)
```

Creates an instance of `ad_ndarray`. This is a low-level method of instantiating an array, it should rather be constructed using factory functions (“empty”, “zeros”, “open”, ...)

**Parameters**

- **shape** (*ints*) – shape of array
- **dtype** – datatype of a single element
- **distaxis** (*int*) – distributed axis
- **target\_offsets** (*ints*) – list of indices marking the offset along the distributed axis of each local array.
- **target\_ranks** (*ints*) – list of *engine* ranks holding the local arrays.
- **no\_allocation** (*bool*) – if `True` no instance of `pyDive.arrays.ad_ndarray.ad_ndarray_local` will be created on engine. Useful for manual instantiation of the local array.
- **kwargs** – additional keyword arguments are forwarded to the constructor of the local array.

**load()**

Load array from file into main memory of all engines in parallel.

**Returns** `pyDive.ndarray` instance

`pyDive.arrays.ad_ndarray.open(filename, datapath, distaxis=0)`

Create a `pyDive.adios.ad_ndarray` instance respectively a structure of `pyDive.adios.ad_ndarray` instances from file.

**Parameters**

- **filename** – name of adios file.
- **datapath** – path within adios file to a single variable or a group of variables.
- **int** (*distaxis*) – distributed axis

**Returns** `pyDive.adios.ad_ndarray` instance

`pyDive.arrays.ad_ndarray.open_variable(filename, variable_path, distaxis=0)`

Create a `pyDive.adios.ad_ndarray` instance from file.

**Parameters**

- **filename** – name of adios file.
- **variable\_path** – path within adios file to a single variable.
- **int** (*distaxis*) – distributed axis

**Returns** `pyDive.adios.ad_ndarray` instance

### 3.1.4 pyDive.cloned\_ndarray package

#### Submodules

##### pyDive.cloned\_ndarray.cloned\_ndarray module

```
class pyDive.cloned_ndarray.cloned_ndarray(shape, dtype=<type  
                                             'float'>, target_ranks=tar-  
                                             get_ranks='all',  
                                             no_allocation=False)
```

Represents a multidimensional, homogenous array of fixed-size elements which is cloned on the cluster

nodes. *Cloned* means that every participating *engine* holds an independent, local numpy-array of the user-defined shape. The user can then do e.g. some manual stuff on the local arrays or some computation with `pyDive.algorithm` on them.

Note that there exists no ‘original’ array as the name might suggest but something like that can be generated by `merge()`.

`__init__(shape, dtype=<type ‘float’>, target_ranks='all', no_allocation=False)`

Creates an `pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray` instance. This is a low-level method for instantiating a cloned\_array. Cloned arrays should be constructed using ‘empty’, ‘zeros’ or ‘empty\_targets\_like’ (see `pyDive.cloned_ndarray.factories`).

#### Parameters

- **shape** (*ints*) – size of the array on each axis
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **target\_ranks** (*ints*) – list of *engine*-ids that share this array. Or ‘all’ for all engines.
- **no\_allocation** (*bool*) – if `True` no actual memory, i.e. *numpy-array*, will be allocated on *engine*. Useful when you want to assign an existing numpy array manually.

`merge(op)`

Merge all local arrays in a pair-wise operation into a single numpy-array.

**Parameters** `op` – Merging operation. Expects two numpy-arrays and returns one.

**Returns** merged numpy-array.

`sum()`

Add up all local arrays.

**Returns** numpy-array.

## pyDive.cloned\_ndarray.factories module

This module holds high-level functions for instantiating `pyDive.cloned_ndarrays`.

`pyDive.cloned_ndarray.factories.empty(shape, dtype=<type ‘float’>)`

Return a new *pyDive.cloned\_ndarray package* utilizing all engines without initializing elements.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.empty_engines_like(shape, dtype, a)`

Return a new `pyDive.cloned_ndarray` utilizing the same engines *a* does without initializing elements.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **a** – *pyDive.arrays.ndarray module*

`pyDive.cloned_ndarray.factories.hollow(shape, dtype=<type ‘float’>)`

Return a new *pyDive.cloned\_ndarray package* utilizing all engines without allocating a local *numpy-array*.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.hollow_engines_like(shape, dtype, a)`

Return a new `pyDive.cloned_ndarray` utilizing the same engines *a* does without allocating a local *numpy*-array.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **a** – *pyDive.arrays.ndarray module*

`pyDive.cloned_ndarray.factories.ones(shape, dtype=<type 'float'>)`

Return a new *pyDive.cloned\_ndarray package* utilizing all engines filled with ones.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.zeros(shape, dtype=<type 'float'>)`

Return a new *pyDive.cloned\_ndarray package* utilizing all engines filled with zeros.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.zeros_engines_like(shape, dtype, a)`

Return a new *pyDive.cloned\_ndarray package* utilizing the same engines *a* does filled with zeros.

#### Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **a** – *pyDive.arrays.ndarray module*

## 3.2 Modules

---

**Note:** All functions of these modules are also directly accessible from the `pyDive` module.

---

### 3.2.1 pyDive.arrayOfStructs module

The *arrayOfStructs* module addresses the common problem when dealing with structured data: While the user likes an array-of-structures layout the machine prefers a structure-of-arrays. In `pyDive` the method of choice is a *virtual array-of-structures-object*. It holds array-like attributes such as `shape` and `dtype` and allows for slicing but is operating on a structure-of-arrays internally.

Example:

```
...
treeOfArrays = {"FieldE" :
                {"x" : fielde_x,
                 "y" : fielde_y,
                 "z" : fielde_z},
                "FieldB" :
                {"x" : fieldb_x,
                 "y" : fieldb_y,
                 "z" : fieldb_z}
               }

fields = pyDive.arrayOfStructs(treeOfArrays)

half = fields[:,2]["FieldE/x"]
# equivalent to
half = fields["FieldE/x"][:,2]
# equivalent to
half = fields["FieldE"]["x"][:,2]
# equivalent to
half = fields["FieldE"][:,2]["x"]

# equivalent to
half = fields.FieldE.x[:,2]
```

The example shows that in fact *fields* can be treated as an array-of-structures **or** a structure-of-arrays depending on what is more appropriate.

The goal is to make the virtual *array-of-structs*-object look like a real array. Therefore every method call or operation is forwarded to the individual arrays.:

```
new_field = fields.FieldE.astype(np.int) + fields.FieldB.astype(np.float)
```

Here the forwarded method calls are `astype` and `__add__`.

`pyDive.arrayOfStructs.arrayOfStructs(structOfArrays)`

Convert a *structure-of-arrays* into a virtual *array-of-structures*.

**Parameters** `structOfArrays` – tree-like dictionary of arrays.

**Raises**

- **AssertionError** – if the *arrays-types* do not match. Datatypes may differ.
- **AssertionError** – if the shapes do not match.

**Returns** Custom object representing a virtual array whose elements have the same tree-like structure as *structOfArrays*.

## 3.2.2 pyDive.algorithm module

`pyDive.algorithm.map(f, *arrays, **kwargs)`

Applies *f* on *engine* on local arrays related to *arrays*. Example:

```
cluster_array = pyDive.ones(shape=[100], distaxis=0)

cluster_array *= 2.0
# equivalent to
pyDive.map(lambda a: a *= 2.0, cluster_array) # a is the local numpy-array of *cluster_array*
```

Or, as a decorator:

```
@pyDive.map
def twice(a):
    a *= 2.0

twice(cluster_array)
```

### Parameters

- **f** (*callable*) – function to be called on *engine*. Has to accept *numpy*-arrays and *kwargs*
- **arrays** – list of arrays including *pyDive.ndarrays*, *pyDive.h5\_ndarrays* or *pyDive.cloned\_ndarrays*
- **kwargs** – user-specified keyword arguments passed to *f*

### Raises

- **AssertionError** – if the *shapes* of *pyDive.ndarrays* and *pyDive.h5\_ndarrays* do not match
- **AssertionError** – if the *distaxis* attributes of *pyDive.ndarrays* and *pyDive.h5\_ndarrays* do not match

### Notes:

- If the hdf5 data exceeds the memory limit (currently 25% of the combined main memory of all cluster nodes) the data will be read block-wise so that a block fits into memory.
- *map* chooses the list of *engines* from the **first** element of *arrays*. On these engines *f* is called. If the first array is a *pyDive.h5\_ndarray* all engines will be used.
- *map* is not writing data back to a *pyDive.h5\_ndarray* yet.
- *map* does not equalize the element distribution of *pyDive.ndarrays* before execution.

`pyDive.algorithm.mapReduce` (*map\_func*, *reduce\_op*, \**arrays*, \*\**kwargs*)

Applies *map\_func* on *engine* on local arrays related to *arrays* and reduces its result in a tree-like fashion over all axes. Example:

```
cluster_array = pyDive.ones(shape=[100], distaxis=0)

s = pyDive.mapReduce(lambda a: a**2, np.add, cluster_array) # a is the local numpy-array of *clu
assert s == 100
```

### Parameters

- **f** (*callable*) – function to be called on *engine*. Has to accept *numpy*-arrays and *kwargs*
- **reduce\_op** (*numpy-ufunc*) – reduce operation, e.g. *numpy.add*.
- **arrays** – list of arrays including *pyDive.ndarrays*, *pyDive.h5\_ndarrays* or *pyDive.cloned\_ndarrays*
- **kwargs** – user-specified keyword arguments passed to *f*

### Raises

- **AssertionError** – if the *shapes* of *pyDive.ndarrays* and *pyDive.h5\_ndarrays* do not match
- **AssertionError** – if the *distaxis* attributes of *pyDive.ndarrays* and *pyDive.h5\_ndarrays* do not match



**Notes:**

- If the hdf5 data exceeds the memory limit (currently 25% of the combined main memory of all cluster nodes) the data will be read block-wise so that a block fits into memory.
- *mapReduce* chooses the list of *engines* from the **first** element of *arrays*. On these engines the mapReduce will be executed. If the first array is a *pyDive.h5\_ndarray* all engines will be used.
- *mapReduce* is not writing data back to a *pyDive.h5\_ndarray* yet.
- *mapReduce* does not equalize the element distribution of *pyDive.ndarrays* before execution.

`pyDive.algorithm.reduce(array, op)`

Perform a tree-like reduction over all axes of *array*.

**Parameters**

- **array** – *pyDive.ndarray*, *pyDive.h5\_ndarray* or *pyDive.cloned\_ndarray* to be reduced
- **op** (*numpy-ufunc*) – reduce operation, e.g. *numpy.add*.

If the hdf5 data exceeds the memory limit (currently 25% of the combined main memory of all cluster nodes) the data will be read block-wise so that a block fits into memory.

### 3.2.3 pyDive.fragment module

`pyDive.fragment.fragment(*arrays, **kwargs)`

Create fragments of *arrays* so that each fragment will fit into the combined main memory of all engines when calling `load()`. The fragmentation is done by array slicing along the distributed axis. The edge size of the fragments is a power of two except for the last fragment.

**Parameters**

- **array** – distributed arrays (e.g. *pyDive.ndarray*, *pyDive.h5\_ndarray*, ...)
- **kwargs** – optional keyword arguments are: `memory_limit` and `offset`.
- **memory\_limit** (*float*) – fraction of the combined main memory of all engines reserved for fragmentation. Defaults to 0.25.
- **offset** (*bool*) – If `True` the returned tuple is extended by the fragments' offset (along the distributed axis). Defaults to `False`.

**Raises**

- **AssertionError** – If not all arrays have the same shape.
- **AssertionError** – If not all arrays are distributed along the same axis.

**Returns** generator object (list) of tuples. Each tuple consists of one fragment for each array in *arrays*.

Note that *arrays* may contain an arbitrary number of distributed arrays of any type. While the fragments' size is solely calculated based on the memory consumption of arrays that store their elements on hard disk (see *hdd\_arraytypes*), the fragmentation itself is applied on all arrays in the same way.

Example:

```
big_h5_array = pyDive.h5.open("monster.h5", "/")
# big_h5_array.load() # crash

for h5_array, offset in pyDive.fragment(big_h5_array, offset=True):
```

```
a = h5_array.load() # no crash
print "This fragment's offset is", offset, "on axis:", a.distaxis
```

`pyDive.fragment.hdd_arraytypes = (<class 'pyDive.distribution.single_axis.h5_ndarray'>, <class 'pyDive.distribution.single_axis.h5_ndarray'>)`  
list of array types that store their elements on hard disk

### 3.2.4 pyDive.mappings module

If `numba` is installed the particle shape functions will be compiled which gives an appreciable speedup.

**class** `pyDive.mappings.CIC`  
Cloud-in-Cell

**class** `pyDive.mappings.NGP`  
Nearest-Grid-Point

`pyDive.mappings.mesh2particles(mesh, particles_pos, shape_function=<class 'pyDive.mappings.CIC'>)`

Map mesh values to particles according to a particle shape function.

#### Parameters

- **mesh** (*array-like*) – n-dimensional array. Dimension of *mesh* has to be greater or equal to the number of particle position components.
- **particles\_pos** (*(N, d)*) – ‘d’-dim tuples for ‘N’ particle positions. The positions can be float32 or float64 and must be within the shape of *mesh*.
- **shape\_function** (*callable, optional*) – Callable object returning the particle assignment value for a given param ‘x’. Has to provide a ‘support’ float attribute which defines the width of the non-zero area. Defaults to cloud-in-cell.

**Returns** Mapped mesh values for each particle.

#### Notes:

- The particle shape function is not evaluated outside the mesh.

`pyDive.mappings.particles2mesh(mesh, particles, particles_pos, shape_function=<class 'pyDive.mappings.CIC'>)`

Map particle values to mesh according to a particle shape function. Particle values are added to the mesh.

#### Parameters

- **mesh** (*array-like*) – n-dimensional array. Dimension of *mesh* has to be greater or equal to the number of particle position components.
- **particles** (*array\_like (1 dim)*) – particle data. `len(particles)` has to be the same as `len(particles_pos)`
- **particles\_pos** (*(N, d)*) – ‘d’-dim tuples for ‘N’ particle positions. The positions can be float32 or float64 and must be within the shape of *mesh*.
- **shape\_function** (*callable, optional*) – Callable object returning the particle assignment value for a given param ‘x’. Has to provide a ‘support’ float attribute which defines the width of the non-zero area. Defaults to cloud-in-cell.

**Returns** *mesh*

#### Notes:

- The particle shape function is not evaluated outside the mesh.

### 3.2.5 pyDive.picongpu module

This module holds convenient functions for those who use pyDive together with `picongpu`.

`pyDive.picongpu.getSteps(folder_path)`

Returns a list of all timesteps in *folder\_path*.

`pyDive.picongpu.loadAllSteps(folder_path, data_path, distaxis=0)`

Python generator object looping hdf5-data of all timesteps found in *folder\_path*.

This generator doesn't read or write any data elements from hdf5 but returns dataset-handles covered by *pyDive.h5\_ndarray* objects.

All datasets within *data\_path* must have the same shape.

#### Parameters

- **folder\_path** (*str*) – Path to the folder containing the hdf5-files
- **data\_path** (*str*) – Relative path starting from “/data/<timestep>” within hdf5-file to the dataset or group of datasets
- **distaxis** (*int*) – axis on which datasets are distributed over when once loaded into memory.

**Returns** tuple of timestep and a *pyDive.h5\_ndarray* or a structure of *pyDive.h5\_ndarrays* (*pyDive.arrayOfStructs*). Ordering is done by timestep.

#### Notes:

- If the dataset has a ‘**sim\_unit**’ attribute its value is stored in *h5array.unit*.

`pyDive.picongpu.loadStep(step, folder_path, data_path, distaxis=0)`

Load hdf5-data from a single timestep found in *folder\_path*.

All datasets within *data\_path* must have the same shape.

#### Parameters

- **step** (*int*) – timestep
- **folder\_path** (*str*) – Path to the folder containing the hdf5-files
- **data\_path** (*str*) – Relative path starting from “/data/<timestep>” within hdf5-file to the dataset or group of datasets
- **distaxis** (*int*) – axis on which datasets are distributed over when once loaded into memory.

**Returns** *pyDive.h5\_ndarray* or a structure of *pyDive.h5\_ndarrays* (*pyDive.arrayOfStructs*).

#### Notes:

- If the dataset has a ‘**sim\_unit**’ attribute its value is stored in *h5array.unit*.

`pyDive.picongpu.loadSteps(steps, folder_path, data_path, distaxis=0)`

Python generator object looping all hdf5-data found in *folder\_path* from timesteps appearing in *steps*.

This generator doesn't read or write any data elements from hdf5 but returns dataset-handles covered by *pyDive.h5\_ndarray* objects.

All datasets within *data\_path* must have the same shape.

#### Parameters

- **steps** (*ints*) – list of timesteps to loop
- **folder\_path** (*str*) – Path to the folder containing the hdf5-files
- **data\_path** (*str*) – Relative path starting from “/data/<timestep>” within hdf5-file to the dataset or group of datasets
- **distaxis** (*int*) – axis on which datasets are distributed over when once loaded into memory.

**Returns** tuple of timestep and a `pyDive.h5_ndarray` or a structure of `pyDive.h5_ndarrays` (*`pyDive.arrayOfStructs`*). Ordering is done by timestep.

#### Notes:

- If the dataset has a ‘**sim\_unit**’ attribute its value is stored in `h5array.unit`.

### 3.2.6 pyDive.pyDive module

Make most used functions and modules directly accessible from `pyDive`.

#### Functions:

`abs`  
`absolute`  
`add`  
`arccos`  
`arccosh`  
`arcsin`  
`arcsinh`  
`arctan`  
`arctan2`  
`arctanh`  
*`array`*  
*`arrayOfStructs`*  
`bitwise_and`  
`bitwise_not`  
`bitwise_or`  
`bitwise_xor`  
`ceil`  
`conj`  
`conjugate`  
`copysign`  
`cos`

cosh  
deg2rad  
degrees  
divide  
empty  
empty\_like  
equal  
exp  
exp2  
expm1  
fabs  
floor  
floor\_divide  
fmax  
fmin  
fmod  
*fragment*  
frexp  
greater  
greater\_equal  
*hollow*  
*hollow\_like*  
hypot  
init  
invert  
isfinite  
isinf  
isnan  
ldexp  
left\_shift  
less  
less\_equal  
log  
log10  
log1p  
log2

logaddexp  
logaddexp2  
logical\_and  
logical\_not  
logical\_or  
logical\_xor  
*map*  
*mapReduce*  
maximum  
*mesh2particles*  
minimum  
mod  
modf  
multiply  
ndarray  
negative  
nextafter  
not\_equal  
ones  
ones\_like  
*particles2mesh*  
power  
rad2deg  
radians  
reciprocal  
*reduce*  
remainder  
right\_shift  
rint  
sign  
signbit  
sin  
sinh  
spacing  
sqrt  
square

subtract  
tan  
tanh  
true\_divide  
trunc  
zeros  
zeros\_like

**Modules:**

IPParallelClient  
*adios*  
*algorithm*  
arrays  
*cloned*  
cloned\_ndarray  
*h5*  
*mappings*  
*picongpu*





## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`

**engine** The cluster nodes of *IPython.parallel* are called *engines*. Sometimes they are also called *targets*. They are the workers of pyDive performing all the computation and file i/o and they hold the actual array-memory. From the user perspective you don't to deal with them directly.



**p**

- `pyDive.algorithm`, [19](#)
- `pyDive.arrayOfStructs`, [18](#)
- `pyDive.arrays.ad_ndarray`, [16](#)
- `pyDive.arrays.h5_ndarray`, [15](#)
- `pyDive.arrays.ndarray`, [12](#)
- `pyDive.cloned_ndarray.factories`, [17](#)
- `pyDive.fragment`, [21](#)
- `pyDive.mappings`, [22](#)
- `pyDive.picongpu`, [23](#)
- `pyDive.pyDive`, [24](#)



## Symbols

`__init__()` (pyDive.arrays.ad\_ndarray.ad\_ndarray method), 15  
`__init__()` (pyDive.arrays.h5\_ndarray.h5\_ndarray method), 14  
`__init__()` (pyDive.cloned\_ndarray.cloned\_ndarray.cloned\_ndarray method), 17  
`__init__()` (pyDive.ndarray method), 11

## A

`ad_ndarray` (class in pyDive.arrays.ad\_ndarray), 15  
`array()` (in module pyDive.arrays.ndarray), 12  
`arrayOfStructs()` (in module pyDive.arrayOfStructs), 19

## C

`CIC` (class in pyDive.mappings), 22  
`cloned_ndarray` (class in pyDive.cloned\_ndarray.cloned\_ndarray), 16  
`copy()` (pyDive.ndarray method), 12

## D

`dist_like()` (pyDive.ndarray method), 12

## E

`empty()` (in module pyDive.arrays.ndarray), 12  
`empty()` (in module pyDive.cloned\_ndarray.factories), 17  
`empty_engines_like()` (in module pyDive.cloned\_ndarray.factories), 17  
`empty_like()` (in module pyDive.arrays.ndarray), 12  
`engine`, 29

## F

`fragment()` (in module pyDive.fragment), 21

## G

`gather()` (pyDive.ndarray method), 12  
`getSteps()` (in module pyDive.picongpu), 23

## H

`h5_ndarray` (class in pyDive.arrays.h5\_ndarray), 14  
`hdd_arraytypes` (in module pyDive.fragment), 22

`hollow()` (in module pyDive.arrays.ndarray), 13  
`hollow()` (in module pyDive.cloned\_ndarray.factories), 17  
`hollow_engines_like()` (in module pyDive.cloned\_ndarray.factories), 18  
`hollow_like()` (in module pyDive.arrays.ndarray), 13

## L

`load()` (pyDive.arrays.ad\_ndarray.ad\_ndarray method), 16  
`load()` (pyDive.arrays.h5\_ndarray.h5\_ndarray method), 14  
`loadAllSteps()` (in module pyDive.picongpu), 23  
`loadStep()` (in module pyDive.picongpu), 23  
`loadSteps()` (in module pyDive.picongpu), 23

## M

`map()` (in module pyDive.algorithm), 19  
`mapReduce()` (in module pyDive.algorithm), 20  
`merge()` (pyDive.cloned\_ndarray.cloned\_ndarray.cloned\_ndarray method), 17  
`mesh2particles()` (in module pyDive.mappings), 22

## N

`ndarray` (class in pyDive), 11  
`NGP` (class in pyDive.mappings), 22

## O

`ones()` (in module pyDive.arrays.ndarray), 13  
`ones()` (in module pyDive.cloned\_ndarray.factories), 18  
`ones_like()` (in module pyDive.arrays.ndarray), 13  
`open()` (in module pyDive.arrays.ad\_ndarray), 16  
`open()` (in module pyDive.arrays.h5\_ndarray), 15  
`open_dset()` (in module pyDive.arrays.h5\_ndarray), 15  
`open_variable()` (in module pyDive.arrays.ad\_ndarray), 16

## P

`particles2mesh()` (in module pyDive.mappings), 22  
`pyDive.algorithm` (module), 19  
`pyDive.arrayOfStructs` (module), 18  
`pyDive.arrays.ad_ndarray` (module), 16  
`pyDive.arrays.h5_ndarray` (module), 15  
`pyDive.arrays.ndarray` (module), 12

pyDive.cloned\_ndarray.factories (module), [17](#)  
pyDive.fragment (module), [21](#)  
pyDive.mappings (module), [22](#)  
pyDive.picongpu (module), [23](#)  
pyDive.pyDive (module), [24](#)

## R

reduce() (in module pyDive.algorithm), [21](#)

## S

sum() (pyDive.cloned\_ndarray.cloned\_ndarray.cloned\_ndarray  
method), [17](#)

## Z

zeros() (in module pyDive.arrays.ndarray), [13](#)  
zeros() (in module pyDive.cloned\_ndarray.factories), [18](#)  
zeros\_engines\_like() (in module py-  
Dive.cloned\_ndarray.factories), [18](#)  
zeros\_like() (in module pyDive.arrays.ndarray), [13](#)