
pyDive Documentation

Release

Author

August 15, 2014

CONTENTS

1	Getting started	3
1.1	Quickstart	3
1.2	Setup an IPython.parallel cluster configuration	3
1.3	Overview	4
2	Tutorials	5
2.1	Example 1: Total field energy	5
2.2	Example 2: Particle density field (picongpu)	6
3	Reference	9
3.1	Packages	9
3.2	Modules	14
4	Indices and tables	21
	Python Module Index	23
	Index	25

Contents:

GETTING STARTED

1.1 Quickstart

pyDive is built on top of *IPython.parallel*, *numpy*, *mpi4py* and *h5py*. Running `python setup.py install` will install pyDive with these and other required packages from *requirements.txt*.

Basic code example:

```
import pyDive
pyDive.init()

arrayA = pyDive.ones([1000, 1000, 1000], distaxis=0)
arrayB = pyDive.zeros_like(arrayA)

# do some array operations, + - * / sin cos, ..., slicing, etc...
...

# get numpy-array
result = arrayC.gather()
# plot result
...
```

Before actually running this script there must have been an *IPython.parallel* cluster started (see section below) otherwise *pyDive.init()* fails.

To keep things simple pyDive distributes array-memory only along **one** user-specified axis. This axis is given by the *distaxis* parameter at array instanciating. It should usually be the largest axis in order to have the best surface-to-volume ratio. But keep in mind that during arithmetic operations both arrays have to be distributed along the *same* axis.

Although the array elements are stored on the cluster nodes you have full access through indexing. If you want to have a numpy-array anyway you can call `pyDive.ndarray.ndarray.ndarray.gather()` knowing that your pyDive array has been sliced down to fit into your local machine's memory.

1.2 Setup an IPython.parallel cluster configuration

The first step is to create an *IPython.parallel* profile in MPI-mode: http://ipython.org/ipython-doc/2/parallel/parallel_process.html . The name of this profile is the argument of `pyDive.init()` . It defaults to "mpi". Starting the cluster is then the second and last step:

```
$ ipcluster start -n 4 --profile=mpi
```

1.3 Overview

pyDive knows three kinds of arrays associated to a separate python package respectively:

- `pyDive.ndarray` -> Stores array elements in cluster nodes' memory.
- `pyDive.h5_ndarray` -> Stores array elements in a hdf5-file.
- `pyDive.cloned_ndarray` -> Holds independent copies of one array on cluster nodes.

Among these three packages there are a few modules:

- `pyDive.arrayOfStructs` -> structured datatypes
- `pyDive.algorithm` -> map, reduce, mapReduce
- `pyDive.mappings` -> particle-mesh mappings
- `pyDive.picongpu` -> helper functions for picongpu-users
- `pyDive.pyDive` -> shortcuts for most used functions

2.1 Example 1: Total field energy

Let's suppose we have a hdf5-file containing a 3D array representing an electric field and we want to compute its total energy. This means squaring and summing or in pyDive's words:

```
import pyDive
import numpy as np
pyDive.init()

h5field = pdDive.h5.fromPath("field.h5", "FieldE", distaxis=0)
field = h5field[:] # read out the entire array into cluster memory in parallel

energy_field = field["x"]**2 + field["y"]**2 + field["z"]**2

total_energy = pyDive.reduce(energy_field, np.add)
```

Now what happens if *h5field* is too large to be stored in the cluster's memory? The line `field = h5field[:]` will crash. In this case we want to load the hdf5 data piece by piece. The functions in `pyDive.algorithm` help us doing so:

```
import pyDive
import numpy as np
pyDive.init()

h5field = pdDive.h5.fromPath("field.h5", "FieldE", distaxis=0)

def square_field(npfield):
    return npfield["x"]**2 + npfield["y"]**2 + npfield["z"]**2

total_energy = pyDive.mapReduce(square_field, np.add, h5field)
```

square_field is called on each *engine* where *npfield* is a structure (`pyDive.arrayOfStructs`) of numpy-arrays representing a sub part of the big *h5field*. `pyDive.algorithm.mapReduce()` can be called with an arbitrary number of arrays including `pyDive.ndarrays`, `pyDive.h5_ndarrays` and `pyDive.cloned_ndarrays`. If there are `pyDive.h5_ndarrays` it will check whether they fit into the cluster memory as a whole and loads them piece by piece if not.

Now let's say our dataset is really big and we just want to get a first estimate of the total energy:

```
...
total_energy = pyDive.mapReduce(square_field, np.add, h5field[::10, ::10, ::10]) * 10.0**3
```

This is valid if *h5field[::10, ::10, ::10]* fits into the cluster's memory. So we could also use the very first version:

```
import pyDive
import numpy as np
pyDive.init()

h5field = pdDive.h5.fromPath("field.h5", "FieldE", distaxis=0)
field = h5field[:, :10, :10, :10]

energy_field = field["x"]**2 + field["y"]**2 + field["z"]**2

total_energy = pyDive.reduce(energy_field, np.add)
```

Note that slicing on a `pyDive.h5_ndarray` always means reading or writing from hdf5 to respectively from memory. But if `h5field[:, :10, :10, :10]` doesn't fit we have to apply the slicing somewhere else in fact at the time when instantiating `h5field`:

```
import pyDive
import numpy as np
pyDive.init()

h5field = pdDive.h5.fromPath("field.h5", "FieldE", distaxis=0, window=np.s_[:, :10, :10, :10])

def square_field(npfield):
    return npfield["x"]**2 + npfield["y"]**2 + npfield["z"]**2

total_energy = pyDive.mapReduce(square_field, np.add, h5field) * 10.0**3
```

This way the hdf5 data is sliced without involving file i/o.

If you use `picongpu` here is an example of how to get the total field energy for each timestep:

```
import pyDive
import numpy as np
pyDive.init()

def square_field(npfield):
    return npfield["x"]**2 + npfield["y"]**2 + npfield["z"]**2

for step, h5field in pyDive.picongpu.loadAllSteps("../simOutput", "fields/FieldE", distaxis=0):
    total_energy = pyDive.mapReduce(square_field, np.add, h5field)

    print step, total_energy
```

2.2 Example 2: Particle density field (picongpu)

Given a huge list of particles in a hdf5-file we want to create a 3D density field out of it. We assume that the particle positions are distributed randomly. This means although each engine is loading a separate part of all particles it needs to write to the entire density field. Therefore the density field must have a whole representation on each participating engine. This is the job of `pyDive.cloned_ndarray`.

```
import pyDive
import numpy as np
pyDive.init()

shape = [256, 256, 256]
density = pyDive.cloned.zeros(shape)
```

```

filename = "../simOutput/h5_1000.h5"
globalCellIdx = pyDive.h5.fromPath(filename, "/data/1000/particles/e/globalCellIdx", distaxis=0)
position = pyDive.h5.fromPath(filename, "/data/1000/particles/e/position", distaxis=0)
weighting = pyDive.h5.fromPath(filename, "/data/1000/particles/e/weighting", distaxis=0)

def particles2density(globalCellIdx, pos, weighting, density):
    total_pos_x = globalCellIdx["x"].astype(pos.dtype["x"]) + pos["x"]
    total_pos_y = globalCellIdx["y"].astype(pos.dtype["y"]) + pos["y"]
    total_pos_z = globalCellIdx["z"].astype(pos.dtype["z"]) + pos["z"]

    # convert total_pos_x, total_pos_y and total_pos_z to an (N, 3) shaped array
    total_pos = np.hstack((total_pos_x[:,np.newaxis],
                           total_pos_y[:,np.newaxis],
                           total_pos_z[:,np.newaxis]))

    import pyDive.mappings
    pyDive.mappings.particles2mesh(density, weighting, total_pos, pyDive.mappings.CIC)

pyDive.map(particles2density, globalCellIdx, position, weighting, density)

final_density = density.sum() # add up all local copies

```

Here, as in the example above, *particles2density* is a function executed on the *engines* by `pyDive.algorithm.map()`. All of its arguments are numpy-arrays or structures (`pyDive.arrayOfStructs`) of numpy-arrays.

REFERENCE

3.1 Packages

3.1.1 pyDive.ndarray package

Submodules

pyDive.ndarray.ndarray module

class `pyDive.ndarray.ndarray.ndarray` (*shape*, *distaxis*, *dtype*=<type 'float'>, *idx_ranges*=None, *targets_in_use*=None, *no_allocation*=False)

Represents a cluster-wide, multidimensional, homogenous array of fixed-size elements. *cluster-wide* means that its elements are distributed across *IPython.parallel-engines*. The distribution is done in one dimension along a user-specified axis. The user can optionally specify which engine maps to which index range or leave the default that pursues an uniform distribution across all engines.

The implementation is based on *IPython.parallel* and local *numpy-arrays*. The design goal is to forward every *numpy-array* method onto the cluster-wide level. Currently `pyDive.ndarray.ndarray.ndarray` supports basic arithmetic operations (+ - * / ** //) as well as most of the *numpy-math* functions like sin, cos, abs, sqrt, ... (see `pyDive.ndarray.dist_math`)

Note that array slicing is a cheap operation since no memory is copied. However this can easily lead to the situation where you end up with two arrays of the same size but of distinct element distribution. Therefore call `dist_like()` first before doing any manual stuff on their local *numpy-arrays*.

Every cluster-wide array operation first equalizes the distribution of all involved arrays if necessary.

__init__ (*shape*, *distaxis*, *dtype*=<type 'float'>, *idx_ranges*=None, *targets_in_use*=None, *no_allocation*=False)

Creates an `pyDive.ndarray.ndarray.ndarray` instance. This is a low-level method for instantiating an array. Arrays should be constructed using 'empty', 'zeros' or 'array' (see `pyDive.ndarray.factories`).

Parameters

- **shape** (*tuple of ints*) – size of the array on each axis
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **idx_ranges** (*tuple/list of (int, int)*) – list of (begin, end) pairs indicating the index range the corresponding *engine* (see *targets_in_use*) is associated with
- **targets_in_use** (*tuple/list of ints*) – list of *engine*-ids that share this array.

- **no_allocation** (*bool*) – if `True` no actual memory, i.e. *numpy-array*, will be allocated on *engine*. Useful when you want to assign an existing *numpy* array manually.

Raises ValueError if just *idx_ranges* is given and *targets_in_use* not or vice versa

If *idx_ranges* and *targets_in_use* are both `None` they will be auto-generated so that the memory will be equally distributed across all *engines* at its best. This means that the last engine may get less memory than the others.

copy ()

Returns a hard copy of this array.

dist_like (*other*)

Redistributes a copy of this array (*self*) like *other* and returns the result. Checks whether redistribution is necessary and returns *self* if not.

Redistribution involves inter-engine communication via MPI.

Parameters other (`pyDive.ndarray.ndarray.ndarray`) – target array

Raises

- **AssertionError** – if the shapes of *self* and *other* don't match.
- **AssertionError** – if *self* and *other* are distributed along distinct axes.

Returns new array with the same content as *self* but distributed like *other*. If *self* is already distributed like *other* nothing is done and *self* is returned.

gather ()

Gathers the local *numpy-arrays* from the *engines*, concatenates them and returns the result.

Returns *numpy-array*

pyDive.ndarray.factories module

This module holds high-level functions for instantiating *pyDive.ndarrays*.

`pyDive.ndarray.factories.array` (*array_like*, *distaxis*)

Return a new *pyDive.ndarray* from an array-like object.

Parameters

- **array_like** (*array-like*) – Any object exposing the array interface, e.g. *numpy-array*, python sequence, ...
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*

`pyDive.ndarray.factories.empty` (*shape*, *distaxis*, *dtype*=<type 'float'>)

Return a new *pyDive.ndarray* distributed across all *engines* without initializing elements.

Parameters

- **shape** (*ints*) – shape of the array
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.ndarray.factories.empty_like` (*a*)

Return a new *pyDive.ndarray* with the same shape, distribution and type as *a* without initializing elements.

`pyDive.ndarray.factories.hollow` (*shape*, *distaxis*, *dtype*=<type 'float'>)

Return a new *pyDive.ndarray* distributed across all *engines* without allocating a local *numpy-array*.

Parameters

- **shape** (*ints*) – shape of the array
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.ndarray.factories.hollow_like(a)`

Return a new *pyDive.ndarray* with the same shape, distribution and type as *a* without allocating a local *numpy*-array.

`pyDive.ndarray.factories.ones(shape, distaxis, dtype=<type 'float'>)`

Return a new *pyDive.ndarray* distributed across all *engines* filled with ones.

Parameters

- **shape** (*ints*) – shape of the array
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.ndarray.factories.ones_like(a)`

Return a new *pyDive.ndarray* with the same shape, distribution and type as *a* filled with ones.

`pyDive.ndarray.factories.zeros(shape, distaxis, dtype=<type 'float'>)`

Return a new *pyDive.ndarray* distributed across all *engines* filled with zeros.

Parameters

- **shape** (*ints*) – shape of the array
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.ndarray.factories.zeros_like(a)`

Return a new *pyDive.ndarray* with the same shape, distribution and type as *a* filled with zeros.

pyDive.ndarray.dist_math module

Mathematical functions supporting `pyDive.ndarray.ndarray.ndarray`:

trigonometric:

`sin`, `cos`, `tan`, `arcsin`, `arccos`, `arctan`

hyperbolic:

`sinh`, `cosh`, `tanh`, `arcsinh`, `arccosh`, `arctanh`

rounding:

`around`, `round`, `rint`, `fix`, `floor`, `ceil`, `trunc`

exponential and logarithmic:

`exp`, `expm1`, `exp2`, `log`, `log10`, `log2`, `log1p`

misc:

`abs`, `sqrt`, `maximum`, `minimum`

3.1.2 pyDive.h5_ndarray package

Submodules

pyDive.h5_ndarray.h5_ndarray module

class `pyDive.h5_ndarray.h5_ndarray(h5_filename, dataset_path, distaxis, window=None)`

Represents a single hdf5-dataset like a virtual, cluster-wide array. Data access goes through array slicing where data is written to respectively read from `pyDive.ndarray.ndarray.ndarray` objects in parallel using all engines.

Example:

```
h5_data = pyDive.h5.fromPath(<file_path>, "/data/0/fields/FieldB/x", distaxis=0)
data = h5_data[:] # read the entire dataset into engine-memory
data = data**2
h5_data[:] = data # write everything back
```

__init__ (*h5_filename, dataset_path, distaxis, window=None*)

Creates an `h5_ndarray` instance. By using this method you may only load a single dataset. If you want to load a *structure* of datasets at once see `pyDive.h5_ndarray.factories.fromPath()`.

Parameters

- **h5_filename** (*str*) – Path of the hdf5-file
- **dataset_path** (*str*) – Path to the dataset within the hdf5-file
- **distaxis** (*int*) – axis on which dataset is to be distributed over during data-access
- **window** (list of slice objects (`numpy.s_`)) – This param let you specify a sub-part of the array as a virtual container. Example: `window=np.s_[::,::2]`

pyDive.h5_ndarray.factories module

`pyDive.h5_ndarray.factories.fromPath(h5_filename, datapath, distaxis, window=None)`

Creates a `pyDive.h5_ndarray` or structure of `pyDive.h5_ndarrays` from a hdf5-dataset respectively a hdf5-group.

Parameters

- **h5_filename** (*str*) – hdf5-filename
- **datapath** (*str*) – path within the hdf5-file to a dataset or a group
- **distaxis** (*int*) – axis on which memory is distributed across the *engines*
- **window** (list of slice objects (`numpy.s_`)) – This param let you specify a sub-part of the array as a virtual container. Example: `window=np.s_[::,::2]`

Returns `pyDive.h5_ndarray` or structure of `pyDive.h5_ndarrays` (`pyDive.arrayOfStructs`)

pyDive.h5_ndarray.h5caching module

`pyDive.h5_ndarray.h5caching.fraction_of_av_mem_used = 0.25`

fraction of the available memory per engine used for caching hdf5 files.

3.1.3 pyDive.cloned_ndarray package

Submodules

pyDive.cloned_ndarray.cloned_ndarray module

```
class pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray(shape, dtype=<type
    'float'>, targets_in_use='all', no_allocation=False)
```

Represents a multidimensional, homogenous array of fixed-size elements which is cloned on the cluster nodes. *Cloned* means that every participating *engine* holds an independent, local numpy-array of the user-defined shape. The user can then do e.g. some manual stuff on the local arrays or some computation with `pyDive.algorithm` on them.

Note that there exists no ‘original’ array as the name might suggest but something like that can be generated by `merge()`.

```
__init__(shape, dtype=<type 'float'>, targets_in_use='all', no_allocation=False)
```

Creates an `pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray` instance. This is a low-level method for instantiating a cloned_array. Cloned arrays should be constructed using ‘empty’, ‘zeros’ or ‘empty_targets_like’ (see `pyDive.cloned_ndarray.factories`).

Parameters

- **shape** (*ints*) – size of the array on each axis
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **targets_in_use** (*ints*) – list of *engine*-ids that share this array. Or ‘all’ for all engines.
- **no_allocation** (*bool*) – if `True` no actual memory, i.e. *numpy-array*, will be allocated on *engine*. Useful when you want to assign an existing numpy array manually.

merge (*op*)

Merge all local arrays in a pair-wise operation into a single numpy-array.

Parameters *op* – Merging operation. Expects two numpy-arrays and returns one.

Returns merged numpy-array.

sum ()

Add up all local arrays.

Returns numpy-array.

pyDive.cloned_ndarray.factories module

This module holds high-level functions for instantiating `pyDive.cloned_ndarrays`.

```
pyDive.cloned_ndarray.factories.empty_engines_like(shape, dtype, a)
```

Return a new `pyDive.cloned_ndarray` utilizing the same engines *a* does without initializing elements.

Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **a** – `pyDive.ndarray`

`pyDive.cloned_ndarray.factories.ones(shape, dtype=<type 'float'>)`

Return a new *pyDive.cloned_ndarray* utilizing all engines filled with ones.

Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.zeros(shape, dtype=<type 'float'>)`

Return a new *pyDive.cloned_ndarray* utilizing all engines filled with zeros.

Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value

`pyDive.cloned_ndarray.factories.zeros_engines_like(shape, dtype, a)`

Return a new *pyDive.cloned_ndarray* utilizing the same engines *a* does filled with zeros.

Parameters

- **shape** (*ints*) – shape of the array
- **dtype** (*numpy-dtype*) – datatype of a single data value
- **a** – *pyDive.ndarray*

3.2 Modules

3.2.1 pyDive.arrayOfStructs module

The *arrayOfStructs* module addresses the common problem when dealing with structured data: While the user likes an array-of-structures layout the machine prefers a structure-of-arrays. In pyDive the method of choice is a *virtual array-of-structures-object*. It holds array-like attributes such as shape and dtype and allows for slicing but is operating on a structure-of-arrays internally.

Example:

```
...
treeOfArrays = {"FieldE" :
                {"x" : fielde_x,
                 "y" : fielde_y,
                 "z" : fielde_z},
                "FieldB" :
                {"x" : fieldb_x,
                 "y" : fieldb_y,
                 "z" : fieldb_z}
               }

fields = pyDive.arrayOfStructs(treeOfArrays)

half = fields[:,2]["FieldE/x"]
# equivalent to
half = fields["FieldE/x"][:,2]
# equivalent to
half = fields["FieldE"]["x"][:,2]
# equivalent to
half = fields["FieldE"][:,2]["x"]
```

The example shows that in fact *fields* can be treated as an array-of-structures **or** a structure-of-arrays depending on what is more appropriate.

The goal is to make the virtual *array-of-structs*-object look like a real array even ready for passing to foreign functions that expect a “real” array which at least inherits e.g. from a numpy-array. Therefore the virtual *array-of-structs*-object inherits from the class type of its arrays. This makes it for instance compatible to `pyDive.algorithm`.

`pyDive.arrayOfStructs.arrayOfStructs (structOfArrays)`

Convert a *structure-of-arrays* into a virtual *array-of-structures*.

Parameters `structOfArrays` – tree-like dictionary of arrays.

Raises

- **AssertionError** – if the *arrays-types* do not match. Datatypes may differ.
- **AssertionError** – if the shapes do not match.

Returns Custom object representing a virtual array whose elements have the same tree-like structure as *structOfArrays*.

3.2.2 pyDive.algorithm module

`pyDive.algorithm.map (f, *arrays, **kwargs)`

Calls *f* on *engine* with local numpy-arrays related to *arrays*. Example:

```
cluster_array = pyDive.ones(shape=[100], distaxis=0)

cluster_array *= 2.0
# equivalent to
pyDive.map(lambda a: a *= 2.0, cluster_array) # a is the local numpy-array of *cluster_array*
```

Parameters

- **f** (*callable*) – function to be called on *engine*. Has to accept *numpy-arrays* and *kwargs*
- **arrays** – list of arrays including *pyDive.ndarrays*, *pyDive.h5_ndarrays* or *pyDive.cloned_ndarrays*
- **kwargs** – user-specified keyword arguments passed to *f*

Raises

- **AssertionError** – if the *shapes* of *pyDive.ndarrays* and *pyDive.h5_ndarrays* do not match
- **AssertionError** – if the *distaxis* attributes of *pyDive.ndarrays* and *pyDive.h5_ndarrays* do not match

Notes:

- If the hdf5 data exceeds the memory limit (see `pyDive.h5_ndarray.h5caching.fraction_of_av_mem_us`) the data will be read block-wise so that a block fits into memory.
- *map* chooses the list of *engines* from the **first** element of *arrays*. On these engines *f* is called. If the first array is a *pyDive.h5_ndarray* all engines will be used.
- *map* is not writing data back to a *pyDive.h5_ndarray* yet.
- *map* does not equalize the element distribution of *pyDive.ndarrays* before execution.

`pyDive.algorithm.mapReduce (map_func, reduce_op, *arrays, **kwargs)`

Calls *map_func* on *engine* with local numpy-arrays related to *arrays* and reduces its result. Example:

```
cluster_array = pyDive.ones(shape=[100], distaxis=0)

s = pyDive.mapReduce(lambda a: a**2, np.add, cluster_array) # a is the local numpy-array of *clu
assert s == 100
```

Parameters

- **f** (*callable*) – function to be called on *engine*. Has to accept *numpy*-arrays and *kwargs*
- **reduce_op** (*numpy-ufunc*) – reduce operation, e.g. *numpy.add*.
- **arrays** – list of arrays including *pyDive.ndarrays*, *pyDive.h5_ndarrays* or *pyDive.cloned_ndarrays*
- **kwargs** – user-specified keyword arguments passed to *f*

Raises

- **AssertionError** – if the *shapes* of *pyDive.ndarrays* and *pyDive.h5_ndarrays* do not match
- **AssertionError** – if the *distaxis* attributes of *pyDive.ndarrays* and *pyDive.h5_ndarrays* do not match

Notes:

- If the hdf5 data exceeds the memory limit (see `pyDive.h5_ndarray.h5caching.fraction_of_av_mem_us`) the data will be read block-wise so that a block fits into memory.
- *mapReduce* chooses the list of *engines* from the **first** element of *arrays*. On these engines the *mapReduce* will be executed. If the first array is a *pyDive.h5_ndarray* all engines will be used.
- *mapReduce* is not writing data back to a *pyDive.h5_ndarray* yet.
- *mapReduce* does not equalize the element distribution of *pyDive.ndarrays* before execution.

`pyDive.algorithm.reduce(_array, op)`
Performs a reduction over all axes of *_array*.

Parameters

- **_array** – *pyDive.ndarray*, *pyDive.h5_ndarray* or *pyDive.cloned_ndarray* to be reduced
- **op** (*numpy-ufunc*) – reduce operation, e.g. *numpy.add*.

If the hdf5 data exceeds the memory limit (see `pyDive.h5_ndarray.h5caching.fraction_of_av_mem_used`) the data will be read block-wise so that a block fits into memory.

3.2.3 pyDive.mappings module

If *numba* is installed the particle shape functions will be compiled which gives an appreciable speedup.

class `pyDive.mappings.CIC`
Cloud-in-Cell

class `pyDive.mappings.NGP`
Nearest-Grid-Point

`pyDive.mappings.mesh2particles(mesh, particles_pos, shape_function=<class pyDive.mappings.CIC at 0x2b5085225808>)`
Map mesh values to particles according to a particle shape function.

Parameters

- **mesh** (*array-like*) – n-dimensional array. Dimension of *mesh* has to be greater or equal to the number of particle position components.
- **particles_pos** ((N, d)) – ‘d’-dim tuples for ‘N’ particle positions. The positions can be float32 or float64 and must be within the shape of *mesh*.
- **shape_function** (*callable, optional*) – Callable object returning the particle assignment value for a given param ‘x’. Has to provide a ‘support’ float attribute which defines the width of the non-zero area. Defaults to cloud-in-cell.

Returns Mapped mesh values for each particle.

Notes:

- The particle shape function is not evaluated outside the mesh.

`pyDive.mappings.particles2mesh(mesh, particles, particles_pos, shape_function=<class pyDive.mappings.CIC at 0x2b5085225808>)`

Map particle values to mesh according to a particle shape function. Particle values are added to the mesh.

Parameters

- **mesh** (*array-like*) – n-dimensional array. Dimension of *mesh* has to be greater or equal to the number of particle position components.
- **particles** (*array_like (1 dim)*) – particle data. `len(particles)` has to be the same as `len(particles_pos)`
- **particles_pos** ((N, d)) – ‘d’-dim tuples for ‘N’ particle positions. The positions can be float32 or float64 and must be within the shape of *mesh*.
- **shape_function** (*callable, optional*) – Callable object returning the particle assignment value for a given param ‘x’. Has to provide a ‘support’ float attribute which defines the width of the non-zero area. Defaults to cloud-in-cell.

Returns *mesh*

Notes:

- The particle shape function is not evaluated outside the mesh.

3.2.4 pyDive.picongpu module

This module holds convenient functions for those who use pyDive together with `picongpu`.

`pyDive.picongpu.getSteps(folder_path)`

Returns a list of all timesteps in *folder_path*.

`pyDive.picongpu.loadAllSteps(folder_path, data_path, distaxis, window=None)`

Python generator object looping hdf5-data of all timesteps found in *folder_path*.

This generator doesn’t read or write any data elements from hdf5 but returns dataset-handles covered by *pyDive.h5_ndarray* objects.

All datasets inside *data_path* must have the same shape.

Parameters

- **folder_path** (*str*) – Path of the folder containing the hdf5-files
- **data_path** (*str*) – Relative path starting from “/data/<timestep>/” within hdf5-file to the dataset or group of datasets

- **distaxis** (*int*) – axis on which datasets are distributed over when once loaded into memory.
- **window** (list of slice objects (*numpy.s_*).) – This param let you specify a sub-part of the array as a virtual container. Example: `window=np.s_[::,::,2]`

Returns tuple of timestep and a *pyDive.h5_ndarray* or a structure of *pyDive.h5_ndarrays* (*pyDive.arrayOfStructs*). Ordering is done by timestep.

`pyDive.picongpu.loadSteps` (*steps, folder_path, data_path, distaxis, window=None*)

Python generator object looping all hdf5-data found in *folder_path* from timesteps appearing in *steps*.

This generator doesn't read or write any data elements from hdf5 but returns dataset-handles covered by *pyDive.h5_ndarray* objects.

All datasets inside *data_path* must have the same shape.

Parameters

- **steps** (*ints*) – list of timesteps to loop
- **folder_path** (*str*) – Path of the folder containing the hdf5-files
- **data_path** (*str*) – Relative path starting from “/data/<timestep>/” within hdf5-file to the dataset or group of datasets
- **distaxis** (*int*) – axis on which datasets are distributed over when once loaded into memory.
- **window** (list of slice objects (*numpy.s_*).) – This param let you specify a sub-part of the array as a virtual container. Example: `window=np.s_[::,::,2]`

Returns tuple of timestep and a *pyDive.h5_ndarray* or a structure of *pyDive.h5_ndarrays* (*pyDive.arrayOfStructs*). Ordering is done by timestep.

3.2.5 pyDive.pyDive module

Make most used functions and modules directly accessible from *pyDive*.

Functions:

`array`
`arrayOfStructs`
`empty`
`empty_like`
`hollow`
`hollow_like`
`init`
`map`
`mapReduce`
`mesh2particles`
`ones`
`ones_like`
`particles2mesh`
`reduce`

`zeros`

`zeros_like`

Modules:

`IPParallelClient`

`algorithm`

`cloned`

`cloned_ndarray`

`h5`

`h5_ndarray`

`mappings`

`ndarray`

`picongpu`

INDICES AND TABLES

- *genindex*
- *modindex*
- *search*

engine The cluster nodes of *IPython.parallel* are called *engines*. Sometimes they are also called *targets*. They are the workers of pyDive performing all the computation and file i/o and they hold the actual array-memory. From the user perspective you don't to deal with them directly.

p

`pyDive.algorithm`, 15
`pyDive.arrayOfStructs`, 14
`pyDive.cloned_ndarray.factories`, 13
`pyDive.h5_ndarray.factories`, 12
`pyDive.mappings`, 16
`pyDive.ndarray.dist_math`, 11
`pyDive.ndarray.factories`, 10
`pyDive.picongpu`, 17
`pyDive.pyDive`, 18

Symbols

`__init__()` (pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray method), 13
`__init__()` (pyDive.h5_ndarray.h5_ndarray.h5_ndarray method), 12
`__init__()` (pyDive.ndarray.ndarray.ndarray method), 9

A

`array()` (in module pyDive.ndarray.factories), 10
`arrayOfStructs()` (in module pyDive.arrayOfStructs), 15

C

`CIC` (class in pyDive.mappings), 16
`cloned_ndarray` (class in pyDive.cloned_ndarray.cloned_ndarray), 13
`copy()` (pyDive.ndarray.ndarray.ndarray method), 10

D

`dist_like()` (pyDive.ndarray.ndarray.ndarray method), 10

E

`empty()` (in module pyDive.ndarray.factories), 10
`empty_engines_like()` (in module pyDive.cloned_ndarray.factories), 13
`empty_like()` (in module pyDive.ndarray.factories), 10
`engine`, 21

F

`fraction_of_av_mem_used` (in module pyDive.h5_ndarray.h5caching), 12
`fromPath()` (in module pyDive.h5_ndarray.factories), 12

G

`gather()` (pyDive.ndarray.ndarray.ndarray method), 10
`getSteps()` (in module pyDive.picongpu), 17

H

`h5_ndarray` (class in pyDive.h5_ndarray.h5_ndarray), 12
`hollow()` (in module pyDive.ndarray.factories), 10
`hollow_like()` (in module pyDive.ndarray.factories), 11

L

`loadAllSteps()` (in module pyDive.picongpu), 17
`loadSteps()` (in module pyDive.picongpu), 18

M

`map()` (in module pyDive.algorithm), 15
`mapReduce()` (in module pyDive.algorithm), 15
`merge()` (pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray method), 13
`mesh2particles()` (in module pyDive.mappings), 16

N

`ndarray` (class in pyDive.ndarray.ndarray), 9
`NGP` (class in pyDive.mappings), 16

O

`ones()` (in module pyDive.cloned_ndarray.factories), 13
`ones()` (in module pyDive.ndarray.factories), 11
`ones_like()` (in module pyDive.ndarray.factories), 11

P

`particles2mesh()` (in module pyDive.mappings), 17
`pyDive.algorithm` (module), 15
`pyDive.arrayOfStructs` (module), 14
`pyDive.cloned_ndarray.factories` (module), 13
`pyDive.h5_ndarray.factories` (module), 12
`pyDive.mappings` (module), 16
`pyDive.ndarray.dist_math` (module), 11
`pyDive.ndarray.factories` (module), 10
`pyDive.picongpu` (module), 17
`pyDive.pyDive` (module), 18

R

`reduce()` (in module pyDive.algorithm), 16

S

`sum()` (pyDive.cloned_ndarray.cloned_ndarray.cloned_ndarray method), 13

Z

`zeros()` (in module pyDive.cloned_ndarray.factories), 14

`zeros()` (in module `pyDive.ndarray.factories`), [11](#)
`zeros_engines_like()` (in module `py-`
`Dive.cloned_ndarray.factories`), [14](#)
`zeros_like()` (in module `pyDive.ndarray.factories`), [11](#)